

Skills Summary

- **Data Engineering & Cloud:** Azure Data Factory , Azure Data Lake Storage Gen2 , Azure Databricks, PySpark, Azure Synapse Analytics, Azure Key Vault, ETL / ELT pipelines. Medallion Architecture (Bronze/Silver/Gold), Data Lakehouse, CI/CD pipelines
- **Programming & Data:** Python (Pandas, NumPy, PySpark), SQL (SQL Server, MySQL), Stored Procedures, Views, Data Validation, Data Reconciliation, API & Flat File Ingestion Frameworks
- **AI / Applied ML:** Retrieval-Augmented Generation (RAG), Semantic Embeddings & Vector Similarity Search, Prompt Engineering (grounded responses), Speech-to-Text pipelines (Whisper), Machine Learning fundamentals.
- **Analytics & BI:** Power BI (DAX, Power Query) Apache Superset, Advanced Excel (Pivot Tables, Lookups).

Databricks Certifications

- Data Ingestion with Lakeflow Connect ([certificate link](#))
- Deploy Workloads with Lakeflow Jobs ([certificate link](#))
- Build Data Pipelines with Lakeflow Spark Declarative Pipeline ([certificate link](#))
- DevOps Essentials for Data Engineering ([certificate link](#))

PROFESSIONAL EXPERIENCE

ADAFSA (Dautom Payroll 3 Months Contract) | Data Engineer| July 2026 – October 2026

- Developed and maintained ETL/ELT pipelines using Azure Data Factory, Azure Data Bricks and Azure data services.
 - Performed data profiling, validation, and transformation using SQL and Python to ensure data quality.
 - Collaborated with business and technical teams to improve data governance and reporting.
- Tools used:** Azure Databricks, Azure Data Factory Pyspark, SQL, Python

PwC | Data Engineer – Analytics Platform | Feb 2025 – July 2026

- Built end-to-end Azure data pipelines using ADF and Databricks (PySpark) to migrate compliance data to Azure Cloud.
 - Designed and implemented 100% end-to-end Azure data pipelines using Azure Data Factory (ADF) and Azure Databricks (PySpark) to migrate and process large-scale compliance data from on-premises systems to Azure Cloud.
 - Performed data transformation, cleansing, and modeling using PySpark, SQL, and Python, enabling reliable analytics datasets for governance reporting across 30,000+ employee records
 - Automated reporting workflows using Python, reducing manual effort by 80%; supported analytics for 30,000+ employee records.
- Tools used:** Azure Data Factory, Azure Databricks, Pyspark, SQL, Python, Power BI, Microsoft Excel

Zepto | Data Engineer (Operation Analytics) | Dec 2023 – Dec 2024

- Developed and maintained SQL-based data pipelines to support logistics and operational analytics.
 - Reduced report turnaround time by 40% by standardizing queries and automating frequently requested data extracts.
 - Performed data cleaning, transformation, and validation in SQL, improving data accuracy by 30% and ensuring reliable reporting outputs
 - Integrated MS SQL Server with Power BI, building dashboards to monitor delivery performance, order metrics, and operational efficiency
- Tools used:** MS SQL Server, Power BI, SQL, Python, Pandas, Apache Superset, Google Sheets, Microsoft Excel

Analytics Vidhya | Analytics Engineer | Nov 2022 – Nov 2023

- Built transformation-ready datasets using SQL and Python for dashboard and analytics use cases.
 - Developed & maintained Dashboards & Reports in Power BI & Excel highlighting sales growth of at least 10-20%.
- Tools used:** Apache Superset, Pandas, SQL, Power BI, Python, MS Excel

PROJECTS

Data Warehouse Project | ETL, CSV, Microsoft SQL Server | [Link](#)

- Designed and implemented a 3-layer (Bronze, Silver, Gold) SQL Server Data Warehouse, integrating data from 2 source systems (ERP & CRM) and processing 10,000+ records into a unified analytical model.
- Built end-to-end ETL pipelines including data cleansing, transformation, and normalization; developed a Star Schema with fact and 5+ dimension tables, improving analytical query performance by 30%.
- Developed 10+ SQL-based analytical reports delivering insights on customer behavior, product performance, and sales trends, enabling data-driven decision-making.

Azure End-to-End Data Engineering Pipeline | Azure Data Factory, ADLS Gen2, Databricks, Synapse, Power BI.

- Built an end-to-end Azure data pipeline ingesting data from external sources using Azure Data Factory.
- Stored raw data in Azure Data Lake Storage Gen2, following a raw → curated layer architecture.
- Performed data transformations using Azure Databricks (PySpark).
- Loaded transformed data into Azure Synapse Analytics for analytical querying.
- Implemented pipeline monitoring and alerts for success/failure notifications.

AI-Powered Q&A Chatbot (RAG Framework) | Python | Whisper | Embeddings | LLMs | [Link](#)

- Designed and implemented a **Retrieval-Augmented Generation (RAG)** chatbot to answer course-related questions with **exact video numbers and timestamps**.
 - Built a pipeline converting **videos** → **audio** → **transcriptions** → **timestamped text chunks**
 - Generated semantic embeddings and applied **vector similarity search** to retrieve relevant content
 - Integrated retrieved context into an LLM using **prompt** to produce grounded, accurate responses
-

EDUCATION

Master of Information Technology – University of Mumbai (2022–2024)

Bachelor of Information Technology – University of Mumbai (2019–2022)